

Tree-Based Speculative Decoding with Adaptive Branching for Automated GPU Kernel Generation

Chandra, Arihan, Ryan, Joel
Stanford University

Abstract—We introduce a novel tree-based speculative decoding framework that leverages hierarchical tree structures with adaptive branching to generate optimized GPU kernels. Our method explores multiple candidate token paths concurrently, reducing idle cycles and improving GPU utilization. By recursively expanding nodes, computing cumulative log probabilities, and dynamically pruning low-probability branches, our approach yields verifiable performance gains. Using a dual-model architecture—where a draft model proposes candidate tokens and a target model verifies and scores sequences—we demonstrate up to $2\text{--}3\times$ throughput improvement and a $40\text{--}60\%$ reduction in token generation latency on KernelBench tasks. This paper details our technical approach, implementation, experimental results, and future directions for automated GPU kernel generation.

I. INTRODUCTION

Automated GPU kernel generation is challenging due to the need for both functional correctness and high performance. Traditional autoregressive decoding methods generate code sequentially, often yielding suboptimal results. We propose a tree-based speculative decoding approach that organizes candidate token sequences into a hierarchical tree. This structure enables parallel exploration of multiple code paths and adaptive pruning of unpromising branches. By leveraging inference-time scaling, our method reduces idle GPU cycles, thereby improving throughput while maintaining output quality.

II. TECHNICAL APPROACH: TREE-BASED SPECULATIVE DECODING WITH ADAPTIVE BRANCHING

Our approach structures the decoding process as a tree where each node represents a candidate token sequence. Key innovations include:

- **Optimized GPU Utilization:** Multiple candidate paths are generated concurrently, minimizing idle cycles.
- **Hierarchical Probability Modeling:** Cumulative log probabilities are computed along each branch, guiding optimal path selection.
- **Dynamic Tree Pruning:** Low-probability branches are randomly pruned when a level exceeds 25 nodes, balancing exploration and efficiency.
- **Verifiable Performance Gains:** A verifier checks for both correctness and performance, ensuring that generated kernels meet quality criteria.

Unlike linear speculative decoding, our tree-based approach maintains a branching structure, enabling parallel exploration of potential solutions.

III. IMPLEMENTATION DETAILS

Our system uses a dual-model architecture:

- **Draft Model (deepseek-ai/DeepSeek-R1-Distill-Qwen-14B):** Proposes candidate tokens.
- **Target Model (deepseek-ai/DeepSeek-R1-Distill-Qwen-32B):** Verifies token sequences and computes cumulative probabilities.

We implement the system in Python using NetworkX for tree management and Matplotlib for visualization. Our custom layout compresses vertical spacing and adds small random horizontal offsets so that branches are not perfectly centered. A random pruning mechanism ensures that no level contains more than 25 nodes. The decoding process is bounded by overall time and node count limits to control runtime.

A. Usage Example

```
result = enhanced_speculative_decode(  
    prompt="Prompt",  
    max_tokens=13,  
    max_nodes_per_level=25,  
    max_total_nodes=200,  
    max_time=10  
)  
print("Generated GPU Kernel Code:")  
print(result)
```

B. Requirements and Installation

- Python 3.7+
- NetworkX, Matplotlib, NumPy
- OpenAI API client
- IPython (for interactive notebook display)
- PyDot (recommended for better tree layouts)

IV. EXPERIMENTAL SETUP AND PERFORMANCE BENCHMARKS

We evaluated our approach on KernelBench, a benchmark with 250 tasks of varying difficulty. Our experiments focused on:

- **Level-1 Tasks:** Single-operator GPU kernels.
- **Level-2 Tasks:** Fused operations requiring multi-stage optimization.

We compared our method with:

- 1) Greedy one-shot generation (no iterative search).
- 2) Beam search/multi-sample generation without feedback.

Our metrics include:

- **Correctness:** Percentage of tasks solved correctly.
- **Latency Reduction:** 40–60% decrease in token generation time.
- **Throughput Improvement:** $2\text{--}3\times$ increase in tokens per second.

V. RESULTS & ANALYSIS

Our tree-based speculative decoding method achieved:

- Near 100% solve rate on Level-1 tasks and 96% on Level-2 tasks within a 15-minute budget per problem.
- A median speedup of $1.7\times$ for Level-1 kernels and $1.5\times$ for Level-2 kernels.
- Consistent output quality, with generated kernels matching or exceeding baseline implementations.

Our experiments show that increased inference-time leads to higher solve rates, with diminishing returns beyond a 20-minute budget. These results confirm that our tree-based approach efficiently balances exploration with computational efficiency.

VI. CONCLUSION & FUTURE WORK

We presented a tree-based speculative decoding framework with adaptive branching and dynamic pruning to generate optimized GPU kernels. Our dual-model system leverages parallel exploration and cumulative probability modeling to achieve significant throughput improvements and reduced latency compared to conventional decoding methods. Future work will focus on further optimizing tree expansion, incorporating adaptive temperature scaling, and extending the framework to handle more complex multi-kernel applications. This approach paves the way for autonomous, high-performance GPU kernel generation driven by large language models.